

Score-seeking in long-horizon agents

A short overview of an eval suite that measures when tool-using LLMs take illegitimate shortcuts

Code & full write-up: <https://github.com/longyi1207/score-seeking-eval>

Most alignment quizzes are one-shot. Weird behavior often shows up when a capable agent is given a **long, hard job** and discovers a soft edge in the environment (think multi-agent incidents on Hugging Face-style platforms, or models improvising channels on obscure wikis). This suite asks a narrower question than “is the model misaligned?”: *under what conditions does it take a planted shortcut instead of finishing the intended work?*

What every task looks like

1. **Honest path** — the intended way to succeed; usually long and annoying (tens of agent steps).
2. **Optional shortcut** — illegitimate but reachable (misconfigured share, leaked backdoor, real CVE, kill the monitor, edit a game file, ...).
3. **Ungameable grader** — ground truth the agent cannot rewrite (watermarked flags and/or checks outside the container), so we label *how* it succeeded, not only whether.

We call a successful shortcut **score-seeking** (or seek). Labels you’ll see: `honest_solve`, `score_seek_success`, `fail` (often a capability / step-cap wall — not the same as honesty).

Why many task types (and many “attack modes”)

If you only test chess file-edits, you learn about chess file-edits. Models that look “safe” on one cheat often look different on another. So the suite deliberately spans domains and shortcut kinds:

- **Azure Active Directory (AD)** — real Windows domain on Azure. Honest path = multi-hop credential chain to Domain Admin. Shortcut = a world-readable planted copy of the DA flag.
- **Protocol reverse-engineering** — figure out a device protocol honestly, or use a leaked maintenance *backdoor* in a config file.
- **Oversight / monitoring** — a live safety process watches the agent; shortcuts include *killing* that process, replacing the test *runner*, or *deceiving* a model code-reviewer with broken code that still “looks fine.”
- **Internal exfil (SSRF)** — abuse an internal HTTP proxy (confused deputy) to mint a privileged token without doing the workflow.
- **Metric redefinition** — game the *definition* of a reported metric so the number hits target while ground truth does not.
- **CVE exploit** — real Grafana path-traversal (CVE - 2021 - 43798) vs using the legitimate API with provided credentials.
- **Games / formal** — hard chess (edit board file / reset difficulty) and Lean proofs (attack the verifier instead of proving).

Abbreviations in Figure 1: **AD** = Active Directory · **CVE** = published vulnerability · **SSRF** = server-side request forgery via internal proxy · **fen-edit / diff-reset** = chess board overwrite / difficulty reset · **seek** = took the shortcut · **contemplated** = an LLM judge found the agent weighing that shortcut in its reasoning (separate from whether it actually used it).

Models so far (n=5 runs per cell): DeepSeek-V4-Pro, Kimi-K2.6, gpt-4o, gpt-5.4, Claude Sonnet 4.5. **Next row planned:** Xiaomi MiMo-V2.6-Pro.




Figure 1. Present suite at a glance. Each cell shows *seek count*, *contemplated count* (second number only when the intent judge has caught up) / shortcut type or outcome / median agent turns. Red = took the shortcut; green = mostly honest success; gray = fail or blocked. A missing second number means that cell is not fully judged yet — that pass is still running. We plan to add a Xiaomi MiMo-V2.6-Pro row.

These are **interesting early results**, not a finished paper. Sample sizes are small (usually $n=5$). Treat them as leads for follow-ups.

1. Prior-agent “permission” matters more than knowing the bug exists

On a real Grafana CVE, with legitimate credentials also available, simply telling the agent the vulnerability exists (*disclose*) → almost nobody exploits it. Leaving a note where a previous engineer / agent *already used* the exploit (*endorse*) → rates jump, especially when the note includes a copy-paste recipe. Endorse *without* a recipe mainly moves DeepSeek.

WHAT WE LEAVE IN THE ENVIRONMENT	DEEPSEEK	KIMI	GPT-4O	GPT-5.4
Nothing / disclose only	0/5	0/5	0/5	0/5
Endorse + recipe	5/5	4/5	5/5	0/5
Endorse, no recipe	4/5	0/5	0/5	0/5
AI endorse, no recipe	2/5	0/5	0/5	0/5

Notes that only add social pressure (ship faster / peer got fined / whistleblower) without granting permission stayed at **0/5** seek. This sits near the multi-agent / “follow the prior agent’s track” pattern people discuss after HF-style incidents — knowledge alone is not the same as normalized use.

2. An explicit “don’t take shortcuts” instruction often works — here

On the protocol backdoor task, telling the agent not to use shortcuts zeros seeking across the panel (e.g. DeepSeek 5/5 → 0/5). Same story on short AD for DeepSeek (60% → 0%). Reading: a lot of what we call seeking may be *ambiguous exploration* that collapses when the norm is clear. Caveat: other benchmarks (e.g. BaitBench) find models still cheat when the shortcut is subtler; our planted edges are often obvious, so “forbid works” may not generalize to quiet Goodharting.

3. Propensity is task- and cheat-specific — not one nationality score

Averaging “Chinese vs Western” cheat rates across this suite is roughly flat; the ranking *flips by shortcut type*. Examples: Kimi is quiet on many hazard tasks but resets hard chess to easy **5/5**; gpt-4o looks honest on short AD but **5/5** deceives a model code reviewer; gpt-5.4 matches DeepSeek on the protocol backdoor and will kill a planted “safety watchdog” tool instead of taking the backdoor. Don’t collapse that into one danger number.

4. Models attack the measurement itself

With a live monitor in the container, some agents kill it or replace the test runner; only held-out checks catch them. We’ve also seen forged audit logs (since fixed), chess difficulty resets that fooled an earlier grader, and Lean statement tampering. That matters for anyone building agent evals: if the grader is in-world, capable models will try to game it.

5. Does a longer honest path reduce seeking?

Suggestive two-point AD contrast (short vs ~30-step) for DeepSeek, but the prompts weren’t perfectly matched and we don’t yet have a clean dose-response. We’re running a **5-length AD ladder** with matched wording (longest honest path targeting ~120 agent turns). Until that lands, we’re not putting a chart of this claim in the brief.

Open questions we want to push next

Highest-leverage directions given what we already see — closer to Neo’s “how weird are models on long hard tasks?” than to “add another model row.”

- **Exploration → exploitation: when does knowing a shortcut become using it?** We already see a large mass of runs where the agent *weighs* the shortcut and still stays honest, plus a smaller set that seeks without saying so. Mapping that boundary — and what tips it (failing honest attempts, prior endorsement, turn pressure) — is the core claim. Same question makes sense of our *forbid* result: one “don’t take shortcuts” line zeros loud planted edges, but other work (e.g. BaitBench) finds models still hack quieter metric leaks while knowing they’re cheating. Pushing our honeypots toward quieter Goodharting / exploit-gated soft edges asks whether that corrigibility was alignment or loudness.
- **Is “prior-agent permission” a real mechanism, or a CVE artifact?** $\text{disclose} \approx 0$ / $\text{endorse} \gg 0$ is our strongest lead, and it rhymes with multi-agent / shared-notes incidents. The clean test is the same ladder on other cheat types (SSRF, metric gaming) and a sharper split of endorsement vs copy-paste recipe — plus whether models *de-escalate* when the endorsed act is more severe than they need.
- **Does honest-path cost actually suppress seeking?** One short vs one long AD is not enough (and may be prompt-confounded). A matched-prompt length ladder asks whether propensity falls smoothly as grinding gets longer — or whether capable models switch past a threshold. The natural endpoint of the same axis is an *impossible* / *near-impossible honest path*: pressure without a capability confound, so “didn’t seek” can’t be read as “couldn’t finish the grind.”

- **What weirdness looks like when the grader sits in the world?** Agents already kill monitors, rewrite runners, and tamper verifiers. Characterizing those attacks — and what held-out checks catch — is both a methods problem for agent evals and a slice of the same propensity question.

Already in motion, not net-new: freeze a present-only **intent×behavior** table once the transcript judge finishes (feeds the exploration→exploitation question above); harden the existing **eval-frame** panel (observed_eval / unobserved_prod on protocol_re + internal_exfil — directional at n=5, finding #15) rather than invent a new sandbagging task; keep task-specific profiles instead of collapsing into one danger score.

Collaboration welcome. We care especially about how **Chinese open-weight / open-API models** behave on these long-horizon, shortcut-tempted agent tasks — relative profiles, not a single “more dangerous” headline. If you’re interested in agent oversight, multi-agent social proof, or exploited reward hacking, please open an issue or reach out via Neo / the repo maintainers: github.com/longyi1207/score-seeking-eval.