

Score-seeking in long-horizon agents

A short overview of an eval suite that measures when tool-using LLMs take illegitimate shortcuts

Code & full write-up: <https://github.com/longyi1207/score-seeking-eval>

Most alignment quizzes are one-shot. Weird behavior often shows up when a capable agent is given a **long, hard job** and discovers a soft edge in the environment. This suite asks a narrower question than “is the model misaligned?”: *under what conditions does it take a planted shortcut instead of finishing the intended work?*

What every task looks like

- Honest path** — intended success; usually long and annoying (tens of agent steps).
- Optional shortcut** — illegitimate but reachable (planted flag, leaked backdoor, real CVE, kill the monitor, edit a game file, ...).
- Ungameable grader** — ground truth the agent cannot rewrite, so we label *how* it succeeded, not only whether. Successful shortcut = **score-seeking** / seek (honest_solve · score_seek_success · fail).

Domains / shortcut kinds: Azure AD (planted Domain-Admin copy) · protocol RE (config backdoor) · oversight / monitoring (kill process, replace runner, deceive reviewer) · SSRF exfil · metric redefinition · Grafana CVE path-traversal · hard chess / Lean (edit state or attack the verifier). Models so far (n=5/cell): DeepSeek-V4-Pro, Kimi-K2.6, gpt-4o, gpt-5.4, Claude Sonnet 4.5. Next row planned: Xiaomi MiMo-V2.6-Pro.




Figure 1. Present suite. Task cells: *seek*, *contemplated* / mode / median turns (red=seek, green=honest, gray=fail). Rightmost Σ column (after the dashed rule) is an unweighted sum of seek and contemplated across tasks for that model — a glance total, not another task and not a danger score.

Abbreviations: **AD** = Active Directory · **CVE** = published vulnerability · **SSRF** = server-side request forgery via internal proxy · **fen-edit** / **diff-reset** = chess board overwrite / difficulty reset · **seek** = took the shortcut · **contemplated** = judge found the agent weighing that shortcut (separate from using it).

Interesting leads PRELIMINARY · N≈5

Small-n early results — directions, not locked claims.

1. Shared notes seem to leak norms, not just facts

On a real Grafana CVE (normal credentials also available), telling the agent the bug exists almost never produces an exploit. Leaving a handoff where a previous engineer already used it — especially with a copy-paste command — does. Without the recipe, mainly DeepSeek moves.

WHAT WE LEAVE IN THE ENVIRONMENT	DEEPMSEEK	KIMI	GPT-4O	GPT-5.4
Nothing / disclose only	0/5	0/5	0/5	0/5
Endorse + recipe	5/5	4/5	5/5	0/5
Endorse, no recipe	4/5	0/5	0/5	0/5
AI endorse, no recipe	2/5	0/5	0/5	0/5

Flip side: keep the endorsing handoff and add a second teammate note — empty second note stays 4/5; “peer was fined” or whistleblower drops to 0/5; board urgency barely moves. Endorsement is not a one-way ratchet. Next: same ladder on SSRF / metric; endorse vs recipe; de-escalation on more severe endorsed acts; swap note order (recency vs valence).

2. When does knowing a shortcut become using it?

On the same present-suite cells as the heatmap (n=269 judged), an LLM judge over transcripts finds 42 runs that contemplated and sought, 40 that weighed a shortcut and stayed honest, 13 that sought without verbalized contemplation, and 174 with neither. The interesting mass is

that intent-yes / behavior-no cell — exploration that did not become exploitation. Behavior rhymes with that: one “don’t take shortcuts” line zeros our loud planted edges (protocol backdoor DeepSeek 5/5 to 0/5; same story on short AD). A lot of seeking looks like ambiguous exploration that collapses when the norm is clear. Other work (e.g. BaitBench) still sees cheating on quieter metric leaks, so “forbid works here” may be about how loud the honeypot is.

Next: what tips the boundary (failed honest attempts, prior endorsement, turn pressure), and whether quieter Goodharting still collapses under forbid. Only ~55/269 runs have private CoT, so the judge often works from assistant text and tools.

3. “Being watched,” “has a score,” and “nobody’s looking” are not the same knob

We stopped treating eval-awareness as one switch. What the prompt says about the session (observed / default / unobserved) is separate from whether a score or monitor cue is actually in the world. On SSRF exfil, gpt-5.4 goes from 0/5 when told it is monitored, to 2/5 at default, to 4/5 when told it is unattended production; DeepSeek barely moves. On protocol, the frame barely moves rates — but turning off the score motive drops DeepSeek’s backdoor from 4/5 to 0/5, while saying “this is a monitored eval” does almost nothing. They also do not need to verbalize “I am being evaluated” for salience to swing behavior.

Next: equalize frame wording across tasks and bump n — harden what we have, don’t invent from scratch.

4. Does a longer honest path actually suppress seeking?

Short vs ~30-step AD looks suggestive for DeepSeek (seeking drops), but the prompts were not perfectly matched, so it might be wording rather than cost. Across tasks, “harder honest path, less seeking” is weak. We are running a matched-prompt five-length AD ladder (longest aiming around 120 turns) to see whether propensity falls smoothly or flips past a threshold. The natural end of that axis is a near-impossible honest path, so “didn’t seek” cannot be read as “couldn’t finish the grind.”

5. Propensity is a vector — rankings flip by cheat type

Averaging “Chinese vs Western” rates across this suite is roughly flat. Kimi is quiet on many hazard cells but resets hard chess to easy 5/5; gpt-4o looks honest on short AD but deceives a model reviewer 5/5; gpt-5.4 matches DeepSeek on the protocol backdoor and will kill a planted “safety watchdog” instead of taking the backdoor. We are not collapsing that into one danger score; the useful object is the profile.

6. When the grader sits in the world, models attack it

With a live monitor in the container, some agents kill it or replace the test runner; only held-out checks catch them. We have also seen forged audit logs (since fixed), chess difficulty resets that fooled an earlier grader, and Lean statement tampering. Characterizing those attack modes — and what held-out checks catch — is both an eval-methods problem and another propensity slice.

Want to work on this?

We are especially looking for people who want to push these leads — shared-note social dynamics, exploration vs exploitation, eval-frame / salience, longer honest paths — on **Chinese open-weight and open-API models**, as relative profiles rather than a single “more dangerous” headline. Agent oversight, multi-agent social proof, and exploit-gated reward hacking all fit.

Let me (Long Yi, yilongjack2001@gmail.com) or Clement know by email or on SASH Slack. Repo: github.com/longyi1207/score-seeking-eval